

Data 612 Assignment 1

Introduction

Our recommender system recommends data science books to readers based on their historical ratings. By predicting ratings when a new reader–book pair is encountered, the system can suggest books the reader is likely to enjoy.

Code

```
import numpy as np
import pandas as pd
from sklearn.metrics import mean_squared_error

data = {
    'user': [
        'A', 'A', 'A',
        'B', 'B', 'B', 'B',
        'C', 'C', 'C', 'C',
        'D', 'D', 'D',
        'E', 'E', 'E', 'E'
    ],
    'item': [
        'book1', 'book2', 'book3',
        'book1', 'book3', 'book4', 'book5',
        'book2', 'book3', 'book4', 'book5',
        'book1', 'book2', 'book3',
        'book1', 'book2', 'book4', 'book5'
    ],
    'rating': [
        5, 3, 4,
        4, 2, 1, 3,
        5, 4, 3, 3,
        3, 3, 4,
        2, 5, 4, 3
    ]
}
```

```
ratings_df = pd.DataFrame(data)
print("Original Ratings:")
print(ratings_df)
```

Original Ratings DataFrame:

	user	item	rating
0	A	book1	5
1	A	book2	3
2	A	book3	4

3	B	book1	4
4	B	book3	2
5	B	book4	1
6	B	book5	3
7	C	book2	5
8	C	book3	4
9	C	book4	3
10	C	book5	3
11	D	book1	3
12	D	book2	3
13	D	book3	4
14	E	book1	2
15	E	book2	5
16	E	book4	4
17	E	book5	3

```

user_item_matrix = ratings_df.pivot(index='user', columns='item',
values='rating')
print("\nUser-Item Matrix:")
print(user_item_matrix)

```

```

User-Item Matrix:
item book1 book2 book3 book4 book5
user
A      5.0   3.0   4.0   NaN   NaN
B      4.0   NaN   2.0   1.0   3.0
C      NaN   5.0   4.0   3.0   3.0
D      3.0   3.0   4.0   NaN   NaN
E      2.0   5.0   NaN   4.0   3.0

```

```
np.random.seed(612)
```

```

mask = np.random.rand(len(ratings_df)) < 0.8
train_df = ratings_df[mask].reset_index(drop=True)
test_df = ratings_df[~mask].reset_index(drop=True)

```

```

print("Training Data:")
print(train_df)
print("\nTest Data:")
print(test_df)

```

```

Training Data:
  user  item  rating
0    A book1      5
1    A book2      3
2    B book1      4
3    B book4      1
4    B book5      3
5    C book3      4

```

6	C	book4	3
7	D	book2	3
8	D	book3	4
9	E	book2	5
10	E	book4	4
11	E	book5	3

Test Data:

	user	item	rating
0	A	book3	4
1	B	book3	2
2	C	book2	5
3	C	book5	3
4	D	book1	3
5	E	book1	2

```

global_avg = train_df['rating'].mean()
print(f"\nGlobal Average Rating: {global_avg:.2f}")

train_df['pred_raw'] = global_avg
test_df['pred_raw'] = global_avg

rmse_train_raw = np.sqrt(mean_squared_error(train_df['rating'],
train_df['pred_raw']))
rmse_test_raw = np.sqrt(mean_squared_error(test_df['rating'],
test_df['pred_raw']))
print(f"Raw Average RMSE (Training): {rmse_train_raw:.2f}")
print(f"Raw Average RMSE (Test): {rmse_test_raw:.2f}")

Global Average Rating: 3.50
Raw Average RMSE (Training): 1.04
Raw Average RMSE (Test): 1.12

user_bias = train_df.groupby('user')['rating'].mean() - global_avg
item_bias = train_df.groupby('item')['rating'].mean() - global_avg

print("\nUser Biases:")
print(user_bias)
print("\nItem Biases:")
print(item_bias)

```

User Biases:

user	
A	0.500000
B	-0.833333
C	0.000000
D	0.000000
E	0.500000

Name: rating, dtype: float64

Item Biases:

item

book1 1.000000

book2 0.166667

book3 0.500000

book4 -0.833333

book5 -0.500000

Name: rating, dtype: float64

```
def baseline_predict(row):  
    u_bias = user_bias.get(row['user'], 0)  
    i_bias = item_bias.get(row['item'], 0)  
    return global_avg + u_bias + i_bias
```

```
train_df['pred_baseline'] = train_df.apply(baseline_predict, axis=1)  
test_df['pred_baseline'] = test_df.apply(baseline_predict, axis=1)
```

```
print("\nTraining predictions:")  
print(train_df[['user', 'item', 'rating', 'pred_baseline']])  
print("\nTest predictions:")  
print(test_df[['user', 'item', 'rating', 'pred_baseline']])
```

Training predictions (first few rows):

	user	item	rating	pred_baseline
0	A	book1	5	5.000000
1	A	book2	3	4.166667
2	B	book1	4	3.666667
3	B	book4	1	1.833333
4	B	book5	3	2.166667
5	C	book3	4	4.000000
6	C	book4	3	2.666667
7	D	book2	3	3.666667
8	D	book3	4	4.000000
9	E	book2	5	4.166667
10	E	book4	4	3.166667
11	E	book5	3	3.500000

Test predictions (first few rows):

	user	item	rating	pred_baseline
0	A	book3	4	4.500000
1	B	book3	2	3.166667
2	C	book2	5	3.666667
3	C	book5	3	3.000000
4	D	book1	3	4.500000
5	E	book1	2	5.000000

```
rmse_train_bl = np.sqrt(mean_squared_error(train_df['rating'],
train_df['pred_baseline']))
rmse_test_bl = np.sqrt(mean_squared_error(test_df['rating'],
test_df['pred_baseline']))
print(f"\nBaseline Predictor RMSE (Training): {rmse_train_bl:.2f}")
print(f"Baseline Predictor RMSE (Test): {rmse_test_bl:.2f}")
```

```
Baseline Predictor RMSE (Training): 0.65
Baseline Predictor RMSE (Test): 1.56
```

Conclusion

In conclusion a simple model was built that predicts every rating as the global average. It gives RMSEs of about 1.04 (training) and 1.12 (test).

Incorporating user and item biases improve training performances, RMSE improving to 0.65, this is because it adjusts for known tendencies. However, on the test set, the RMSE increased to 1.56. This means that the biases capture the training data well, but they overfit on the test dataset. However this is to be expected since we are working with a toy dataset with very few data.